

Background and requirements:

The goal in this project is to train two agents controlling rackets to bounce a ball over a net. By taking advantage of the Unity Machine Learning Agents (ML-Agents), which is an open-source Unity plugin that enables games and simulations to serve as environments for training intelligent agents, this project would be straightforward on focusing on the model development and performance improving.

In this environment, two agents control rackets to bounce a ball over a net. If an agent hits the ball over the net, it receives a reward of +0.1. If an agent lets a ball hit the ground or hits the ball out of bounds, it receives a reward of -0.01. Thus, the goal of each agent is to keep the ball in play.

The observation space consists of 8 variables corresponding to the position and velocity of the ball and racket. Each agent receives its own, local observation. Two continuous actions are available, corresponding to movement toward (or away from) the net, and jumping.

The task is episodic, and in order to solve the environment, your agents must get an average score of +0.5 (over 100 consecutive episodes, after taking the maximum over both agents). Specifically,

- After each episode, we add up the rewards that each agent received (without discounting), to get a score for each agent. This yields 2 (potentially different) scores. We then take the maximum of these 2 scores.
- This yields a single score for each episode.

The environment is considered solved, when the average (over 100 episodes) of those scores is at least +0.5.

Learning Algorithm

After briefly reviewing the project requirement, a straightforward idea using the solutions from the project 2 but with MADDPG algorithm (Multi-Agent Deep Deterministic Policy Gradient) looking very promising.

Policy-based, Value-based and Actor-Critic:


RL algorithms using neural networks as function approximators can take approaches with things they are estimating. Firstly, given a state as input they can estimate the expected future rewards for different possible actions. This approach is known as a **value-based method**. However, for tasks that involve continuous actions, a function approximator to estimate the actual policy directly could be applied. This approach is known as a **policy-based method**. By applying the DDPG (Deep

Deterministic Policy Gradient, Continuous Action-space) algorithm as an "Actor-Critic" method is introduced.

MADDPG:

As presented in the paper Multi-Agent Actor-Critic for Mixed Cooperative-Competitive Environments (<https://arxiv.org/abs/1706.02275>), this project adapts the actor-critic method mentioned above to consider action policies of other agents by adopting a centralized training with decentralized execution.

Decentralized: any agent has its private Actor network and takes actions considering only its observation of the environment. **Centralized training:** although each agent has its private Critic network, the critic takes as input the observations of any agent present in the environment as well as their actions in order to ease training while learning the action-value function $Q(s,a)$.




On top of taking advantage of the MADDPG algorithm, this solution also had several adjustments has been done to improve the performance,

1. On top of the Ornstein-Uhlenbeck process, introduce the epsilon parameter for noise decay
2. Fine-tune the theta (speed of mean reversion) and sigma (volatility) parameters for adding noise
3. Applying gradient clipping and normalization
4. Adjusting learning interval, by learn every 20 steps with 10-time multiplication.

After these adjustments, the performance improved and the environment is solved as shown in the plot, the agent is able to receive an average (over 100 episodes) of those scores is at least +0.5.

```
[8]: import matplotlib.pyplot as plt
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(1, len(scores)+1), scores)
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.show()
```



```
Episode 100    Average Score: 0.02range maximum score over the last 10 episodes: 0.00
Episode 200    Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 300    Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 400    Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 500    Average Score: 0.00range maximum score over the last 10 episodes: 0.01
Episode 600    Average Score: 0.00range maximum score over the last 10 episodes: 0.01
Episode 700    Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 800    Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 900    Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 1000   Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 1100   Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 1200   Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 1300   Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 1400   Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 1500   Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 1600   Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 1700   Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 1800   Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 1900   Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 2000   Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 2100   Average Score: 0.00range maximum score over the last 10 episodes: 0.00
Episode 2200   Average Score: 0.03range maximum score over the last 10 episodes: 0.05
Episode 2300   Average Score: 0.01range maximum score over the last 10 episodes: 0.03
Episode 2400   Average Score: 0.01range maximum score over the last 10 episodes: 0.00
Episode 2500   Average Score: 0.06range maximum score over the last 10 episodes: 0.13
Episode 2600   Average Score: 0.07range maximum score over the last 10 episodes: 0.01
Episode 2700   Average Score: 0.09range maximum score over the last 10 episodes: 0.08
Episode 2800   Average Score: 0.22range maximum score over the last 10 episodes: 0.10
Episode 2866   max score: 2.60 average maximum score over the last 10 episodes: 1.46
Environment solved in 2766 episodes!    Average Score: 0.51
```

Model architecture

ACTOR NETWORK

Batch Norm Layer -> Fully Connected Layer -> Leaky Relu(leakiness=0.01) -> Fully Connected Layer -> Leaky Relu(leakiness=0.01) Fully Connected Layer -> Tanh Activation()

CRITIC NETWORK

Batch Norm Layer -> Fully Connected Layer -> Leaky Relu(leakiness=0.01) -> Fully Connected Layer -> Leaky Relu(leakiness=0.01) Fully Connected Layer

Hyperparameters and values

Continuous_control.ipynb Notebook:

- rewards # get the rewards
- next_states # get the resulting states
- dones = env_info.local_done # check whether episodes have finished

ddpg_agent.py:

- BUFFER_SIZE = int(1e6) # replay buffer size
- BATCH_SIZE = 128 # minibatch size
- GAMMA = 0.99 # discount factor
- TAU = 1e-3 # for soft update of target parameters
- LR_ACTOR = 1e-3 # learning rate of the actor
- LR_CRITIC = 1e-3 # learning rate of the critic
- WEIGHT_DECAY = 0. # L2 weight decay
- OU_SIGMA = 0.2 # Ornstein-Uhlenbeck noise parameter
- OU_THETA = 0.15 # Ornstein-Uhlenbeck noise parameter
- EPSILON = 1.0 # explore->exploit noise process added to act step
- EPSILON_DECAY = 1e-6 # decay rate for noise process

model.py:

- state_size (int): State dimension
- action_size (int): Action dimension
- random_seed (int): Random seed
- fc1_units (int): First hidden layer nodes
- fc2_units (int): second hidden layer nodes
- hidden_size(int): Hidden layers node
- leak Leakiness in leaky relu

Future ideas

As this project is done in tight timeline, there still plenty to improve when time allows, that includes:

- **Hyperparameters tuning**, that would give a better performance and solve the environment faster.
- Improve the **noise process algorithms**.
- Training regimen utilizing an **ensemble of policies** for each agent.
- Further investigation of introducing **negative rewards**.